

AI for Software Engineering

June 28, 2025

0.1 # AI for Software Engineering Assignment # Title: Understanding the AI Development Workflow 0

0.2 Part 1: Short Answer Questions (30 pos)

0.2.1 1. Problem Definition (6 ints)

Hypothetical Problem:

Recommending local job opportunities to unemployed youth based on skills and location. *Objectives: **1. Match user skills to relevant job descriptions using NLP. 2. Increase accessibility to local job listings in under-resourced areas. 3. Reduce the job search time by delivering personalized recommendations.** Stakeholders:** - Youth/job seekers in townships and rural areas. - Local employers or small businesses posting jobs.

KPI (Performance Indicator):

Job Match Accuracy — percentage of recommended jobs clicked for by users.

0.2.2 2. Data Collection Preprocessing (8 points)

Data Sources: 1. Local job postings scraped from WhatsApp groups, Facebook pages, or community websites. 2. User-submitted profiles with skills and preferred location.

Potential Bias:

Location bias — over-representation of job posts from urban areas, underrepresenting rural communities.

Preprocessing Steps: 1. Text cleaning — remove stopwords and punctuation from job descriptions. 2. Vectorization — convert text to numerical form using TF-IDF. 3. Handling missing data — fill in missing fields like location using location inference.

•

1 3. Model Development (8 points)

**Model: TF-IDF + Cosine Similarity

Justification: Lightweight and effective for text matching.

Data Split

70% training / 15% validation / 15% test

Hyperparameters to Tune: 1. Max number of features in TFF.. N-gram range (e.g., unigrams and bigrams).

1.0.1 4. valuation & Deployment (8 points)

Evaluation Metrics: 1. Precision — to ensure relevant job mates. 2. Click-through rate (CTR) — tracks user engagement.

Concept Drift: Occurs when job language or user behavior changes over time.

Monitoring Strategy: Track CTRs and periodically retrain the model.

Deployment Challenge:

Scale — handling large amounts of user and job data in real time.

1.1 Part 2: Case duplication (40 points)

1.1.1 Hospital Readmission Risk Prediction

1.1.2 1. Problem Scope (5 points)

Problem Statement:

Predict likelihood of patient readmission within 30 days of discharge.

Objectives: 1. Identify high-risk patients. 2. Allocate post-discharge care efficiently. 3. Reduce hospital costs and improve outcomes. **Stakeholders:** - Hospital management - Medical staff (doctors, discharge planners)

1.1.3 2. Data Strategy (10 points)

Data sources: 1. Electronic Health Records (EHRs) 2. Demographics (e.g., age, gender, past visits)

Ethical Concerns: 1. Patient privacy (HIPAA compliance) 2. Informed consent to use data for AI modeling

Preprocessing Pipeline: - Clean and anonymize data. - Engineer features (e.g., age at length, number of comorbidities) - Normalize numeric values and encode categorical variables.

2 3. Model Development (10 points)

Model: Gradient Boosting Classifier (e.g., XGBoost) **Justification:**

Performs well on structured healthcare data, supports missing values, and is interpretable.

Hypothetical Confusion Matrix:

	Predicted: No	Predicted: Yes
Actual: No**	85	15
Actual: Yes	10	40

Precision = $40 / (40 + 15) = 72.7\%$

Recall = $40 / (40 + 10) = 80\%$

2.0.1 4. Deployment (10 points)

Integration Steps: 1. Wrap model in a REST API (Flask or FastAPI). 2. Connect it to the hospital's EHR system. 3. Display risk predictions in the discharge planning dashboard.

****Compliance Measures:** * Egypt patient data (in transit & at rest). - Restrict model access to authorized personnel. - Follow HIPAA or POPIA regulations.

2.0.2 5. Operation (5 points)

Overfitting Mitigation: - Use Cross-Validation for evaluation. - Apply Regularization (L1/L2) in the model.

2.1 Part 3: Critical Thinking (20 points)

2.1.1 1. Ethics & Bias (10 points)

Impact of Biased Training Data: - May underpredict risk for under-represented groups. - Can result in missed interventions or unfair prioritization. **Strategy:** ** - Balance the dataset using re-sampling. - Perform fairness evaluation across demographic groups. - Use fairness-aware algorithms.

2.1.2 2. Trade-offs (10 points)

Interpretability vs. Accuracy: - Interpretable models (e.g., decision trees) improve trust but may lose accuracy. - High-accuracy models (e.g., deep learning) may lack transparency. - In healthcare, interpretability is prioritized. **Impact of Limited Resources:** - Use lightweight (e.g., logistic regression) - Avoid GPU-heavy approaches. - Consider batch rather than real-time inference.

2.2 ## Part 4: Reflition & Workflow Diagram (10 points)


2.2.1 1. Reflection (5 points)

Most Challenging Stage:

Data collection and preprocessing due to privacy, bias, and quality concerns.

Improvements with Moimesources: - Partner with medical institutions for real-world data.
- Include domain experts in the feature engineering process. - Use federated learning to protect patient privacy.

2.2.2 2. AI Workflow Diagram (Text Version) (5 points)



```
[ ]: ---

Would you like help uploading this to GitHub right now or adding more features_
↳like real CSV data or a simple Streamlit app? | :contentReference[oaicite:
↳0]{index=0}|
```

```
[ ]: computer, office, admin
```

```
[1]: import sklearn
print(sklearn.__version__) # Should show version (e.g., 1.3.0)
```

1.3.0

```
[ ]: # job_recommender.py

import pandas as pd
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
import numpy as np

# -----
# 1. Load Sample Data
# -----

# Sample job descriptions
jobs = pd.DataFrame({
    'job_id': [1, 2, 3, 4, 5],
    'title': ['Admin Assistant', 'Junior Web Developer', 'Data Entry Clerk',
↳'Warehouse Packer', 'Receptionist'],
```

```

        'description': [
            'Looking for an admin assistant with Microsoft Office skills and
↪attention to detail.',
            'Seeking a junior web developer with HTML, CSS, and JavaScript
↪experience.',
            'Data entry clerk needed for capturing info into spreadsheets. Accuracy
↪important.',
            'Warehouse packer needed for sorting goods and packing orders. Physical
↪strength required.',
            'Receptionist needed with good communication and computer literacy.'
        ],
        'location': ['Soweto', 'Johannesburg', 'Tembisa', 'Diepsloot', 'Alexandra']
    })

# -----
# 2. Preprocess and Vectorize
# -----

# Combine relevant fields for matching
jobs['text'] = jobs['title'] + " " + jobs['description']

# Vectorize using TF-IDF
vectorizer = TfidfVectorizer(stop_words='english', max_features=1000)
job_vectors = vectorizer.fit_transform(jobs['text'])

# -----
# 3. User Input and Matching
# -----

def recommend_jobs(user_skills, top_n=3):
    user_vector = vectorizer.transform([user_skills])
    similarities = cosine_similarity(user_vector, job_vectors).flatten()
    top_indices = np.argsort(similarities)[-top_n:][::-1]
    return jobs.iloc[top_indices][['job_id', 'title', 'location',
↪'description']]

# Example usage
if __name__ == "__main__":
    print("=== Youth Job Recommender ===")
    user_input = input("Enter your skills (e.g. computer, office, admin): ")
    results = recommend_jobs(user_input)
    print("\nTop Job Matches:")
    for _, row in results.iterrows():
        print(f>Title: {row['title']} | Location: {row['location']}")
        print(f>Description: {row['description']}")
        print("-" * 40)

```

```
=== Youth Job Recommender ===
```

```
[ ]:
```